

# Lecture 11

## Byzantine Fault Tolerance

ECE 422: Reliable and Secure Systems Design



Instructor: An Ran Chen  
Term: 2024 Winter

# Schedule for today

- Key concepts from last class
- Failure classification
- Halting failures
- Byzantine failure
- The Two Generals Problem
- Next class: Byzantine General problem

# Example on encoding



Suppose  $g(x) = (1+x+x^3)$  for a  $(7,4)$  cyclic code

**Question:** Find the codewords for the following data: 0001, 1001, 0110, 1000

**Solution:** For each entry in the data (e.g., 0001, 1001):

- Step 1: convert it into a data polynomial
- Step 2: solve the codeword polynomial multiplication

Hint: use the circular shifting property to calculate the codeword polynomial instead. It is much faster.

# Example on encoding



Suppose  $g(x) = (1+x+x^3)$  for a (7,4) cyclic code

No shifting, 1 shifting, 3 shifting

**Question:** Find the codewords for the following data: 0001, 1001, 0110, 1000

**Solution:** Using the circular shifting property, compute for  $d(x) \cdot g(x)$

For data = {0001000},  $d(x) \cdot g(x) = 0001000 + 0000100 + 0000001 = 0001101$

# Summary of cyclic code

- Any circular shift of a codeword produces another codeword.
- Code is characterized by its generator polynomial  $g(x)$ , with a degree  $(n-k)$ , where  $n$  = bits in codeword,  $k$  = bits in data.
- All calculations are done in mod 2 arithmetic.
  - Multiplication of polynomial for encoding
  - Division of polynomial for decoding
- Cyclic code detects all single errors and all multiple adjacent error affecting  $(n-k)$  bits or less.
  - It does not correct the error.

# Information redundancy

Information redundancy tolerate faults by adding information to the original data.

- Tolerate faults by means of coding
- Avoid unwanted information changes
- E.g., information loss during data storage or transmission

The term “coding” is used in the context of communication and data storage

# Software fault tolerance

Fault tolerance is the ability for systems to continue functioning as a whole despite the faults

- Example of fault tolerance: while a PDF reader may crash when editing a corrupted PDF file, the PDF reader will restart itself after saving the information

The system contains a **fail-stop failure**.

- Upon a fault, the system stops.
- To tolerate this type of behavior:
  - Solution 1: Restart the system (e.g., resetting the states)
  - Solution 2: Failover to redundancy


# Schedule for today

- Key concepts from last class
- Failure classification
- Halting failures
- Byzantine failure
- The Two Generals Problem
- Next class: Byzantine General problem



# Failure

Failure occurs when the system fails to perform its required function.



A failure is always related to a required function

- Based on the system specifications and functional requirements
- E.g., according to the user requirement, the maximum response time of a server shall be no longer than 15 seconds.
- A failure occurs when the response time exceeds 15 seconds.

# Failure

A failure is an event that occurs at a specific point in time.

- It is not always possible to observe when a failure happens
- E.g., server performance degradation
- E.g., dormant faults in codes

The failure may:

- Come in different forms
- Originate from different kinds of faults
- Have different consequences on the system

# Failure classification

Failure classification describes the way how a system can fail that is perceived by the rest of the system.

- It aims to answer one question: “What kind of failure are we dealing with when considering the whole system into account.”
- It helps understand how to build and design for fault-tolerant systems.

There are four types of failure:

- Timing failure, also known as performance failure
- Omission failure
- Crash failure
- Arbitrary failure, also known as Byzantine failure

# Timing failure

Timing failure happens when the system delivers correct results, but outside the expected time interval.

- Failure lies in the amount of time the system took to execute a task
- Not in the results of the task itself

Example: Delayed response from a server

- The server must respond between 5 milliseconds and 15 seconds
- A response exceeds our expected time interval
- This produces a timing failure

# Omission failure

Omission failure happens when the system never appears to respond.

- Can also be understood as “infinitely late” timing failure

There are two forms of omission failure:

- Send omission failure: fail to send a response
- Receive omission failure: fail to receive a response

Example: No response to incoming requests

- The client sends a request
- The server fails to respond to the request
  - Failed to send messages
  - Failed to receive messages

# Crash failure

Crash failure happens when the system crashes, but is working correctly until it crashes.

- The system experiences omission once and becomes completely unresponsive (i.e., crashes)

Example: Server crashes

- The client sends a request
- The server experiences an omission failure, then fails and stops responding,

# Byzantine failure

Byzantine failure happens when the system sends arbitrary responses at arbitrary times.

- Derived from Byzantine faults
- The system responds with different (erroneous) responses each time, and is inconsistent in what kind of response it delivers

Example: Arbitrary response to the same request

- The client A sends a request
- The server responds with a message A
- The client B sends the same request as client A
- The server arbitrarily responds with a message B


# Examples of failure



Type of failures	Example of server's response
<b>Crash failure</b>	Valid responses until it crashes
<b>Omission failure</b>	No response to incoming requests
<b>Timing failure</b>	Delayed response, outside of a specified time interval
<b>Arbitrary failure</b>	Arbitrary response at arbitrary times



# Difficulties of failures



# Halting failures

Failure classification deal with the situation that we no longer perceive any actions from the system.

- Can we conclude that a system has come to a halt?
- How do we distinguish between crash/omission/timing failure?

There are five categories of halting failures on the reliability of failure detection (from the least to the most severe):

- Fail-stop
- Fail-noisy
- Fail-silent
- Fail-safe
- Fail-arbitrary

# Types of halting failures

Let a *client* attempt to detect that the *server* has failed.

Fail-stop: failures that can be reliably detected.

- Crash failure, but reliably detectable
- Assuming that:
  - Non-faulty communication exists between the client and the server
  - There is a worst-case delay on the response from the server.
- E.g., Server stops and the client can detect this failure.

Fail-noisy: failures that are eventually reliably detected

- Crash failure, but eventually reliably detected
- The client will only eventually come to the correct conclusion that the server has crashed.
- E.g., Server's response time slows down, and eventually fails.

# Types of halting failures

Fail-silent: failures that cannot be distinguished between crash and omission

- Omission or crash failures
- Assuming that:
  - Non-faulty communication exists between the client and the server
  - The client cannot distinguish crash failures from omission failures.
- E.g., No response from the server.

Fail-safe: failures that cannot do any harm.

- Arbitrary, yet benign failures
- The server fails without doing any harm
- E.g., Server slows down, but remains available.

# Types of halting failures

Fail-arbitrary: arbitrary, with malicious failures

- Server may fail in any possible way; failures may be unobservable.
- Arbitrary where a server is producing output that it should never have produced, but which cannot be detected as being incorrect.
- E.g., The server may be working with other servers to produce intentionally wrong answers; it still responds to the client.

# Schedule for today

- Key concepts from last class
- Failure classification
- Halting failures
- **Byzantine failure**
- The Two Generals Problem
- Next class: Byzantine General problem

# Byzantine failure

Byzantine failure: a node/component may fail arbitrary due to:

- Exhausted resources
  - E.g., null response from the server due to resource exhaustion
- Conflicting information from different parts of the system
  - E.g., Malicious attack providing conflicting information
- ... and more

# Byzantine failure

Byzantine failure: a node/component may fail arbitrary due to:

- Exhausted resources
- Conflicting information from different parts of the system

Why would nodes/components fail arbitrarily?

- Software bugs in the code
- Hardware failures
- Malicious attack on the system



# What makes Byzantine failure important?

Byzantine failure provides an abstraction of issues in real-world applications

Example 1: Bitcoin

- Byzantine failure as the results of malicious nodes in the network
- A malicious node can cause failure in the network

Example 2: Aircraft flight control system (e.g., Boeing 777, 787)

- Byzantine failure as the results of potential hardware failures
- Hardware components can fail arbitrarily based on the temperature at which they operate

Solution? Designing Byzantine fault-tolerant systems for reliability

# The Two Generals Problem


The two generals problem is a consensus problem where two Generals (General A1 and General A2) need to coordinate attack on the army B.

- If both Generals attack at the same time, they win the battle.
- If neither Generals attack, they will try to coordinate again the next day.
- But if only one of the two Generals attacks, they lose the battle.

There is only one way to communicate between the two Generals:

- By sending messengers on a high-risk path through enemy lines.

What messages can the generals send to reach a consensus on whether or not to attack?



# The Two Generals Problem

Let General A1 be the one who decides to attack the next morning.

- Send the message to General A2 about the attack.
- Make sure that General A2 got the message.

What kind of message can General A1 send to General A2 to make sure that a consensus is reached?

# Solution: TCP 3-way handshake

**A1**



A1 wants to  
attack

**A2**

*If you respond, I'll attack!*

A1 will attack;  
B wants to attack

Scenario 1: If the  
message was delivered  
to A2, everything is fine.

# Solution: TCP 3-way handshake

**A1**



A1 wants to  
attack

*If you respond, I'll attack!*



**A2**

A1 will attack;  
A2 wants to attack

Scenario 2: If the messenger was captured by B, A1 will not receive any response from A2, they will not attack.

# Solution: TCP 3-way handshake

## A1



## A2

A1 wants to  
attack

*If you respond, I'll attack!*

A1 will attack;  
A2 will attack

*If you respond, I'll attack!*

A1 will attack;  
A2 wants to attack

Scenario 1: If the message  
was delivered back to A1,  
everything is fine.

# Solution: TCP 3-way handshake

## A1



## A2

A1 wants to attack

*If you respond, I'll attack!*

A1 will attack;  
A2 wants to attack

A1 will attack;  
A2 will attack

*If you respond, I'll attack!*



Scenario 2: If the messenger was captured by B, A2 will not receive any response from A1, they will not attack.

# Solution: TCP 3-way handshake

## A1



## A2

A1 wants to attack

*If you respond, I'll attack!*

A1 will attack;  
A2 will attack

*If you respond, I'll attack!*

*We'll attack!*

A1 will attack;  
A2 wants to attack

A1 will attack;  
A2 will attack

Scenario 1: If the message was delivered back to A2, everything is fine. They will both attack.



# Solution: TCP 3-way handshake

## A1



## A2

A1 wants to attack

*If you respond, I'll attack!*

A1 will attack;  
A2 will attack

*If you respond, I'll attack!*

*We'll attack!*



A1 will attack;  
A2 wants to attack

A1 will attack;  
A2 will attack

Scenario 2: If the message was not delivered back to A2, A1 is committed to attack but A2 isn't.

# Solution: TCP 3-way handshake

**A1**



**A2**

A1 wants to  
attack

*If you respond, I'll attack!*

A1 will attack;  
A2 wants to attack

A1 will attack;  
A2 will attack

*If you respond, I'll attack!*

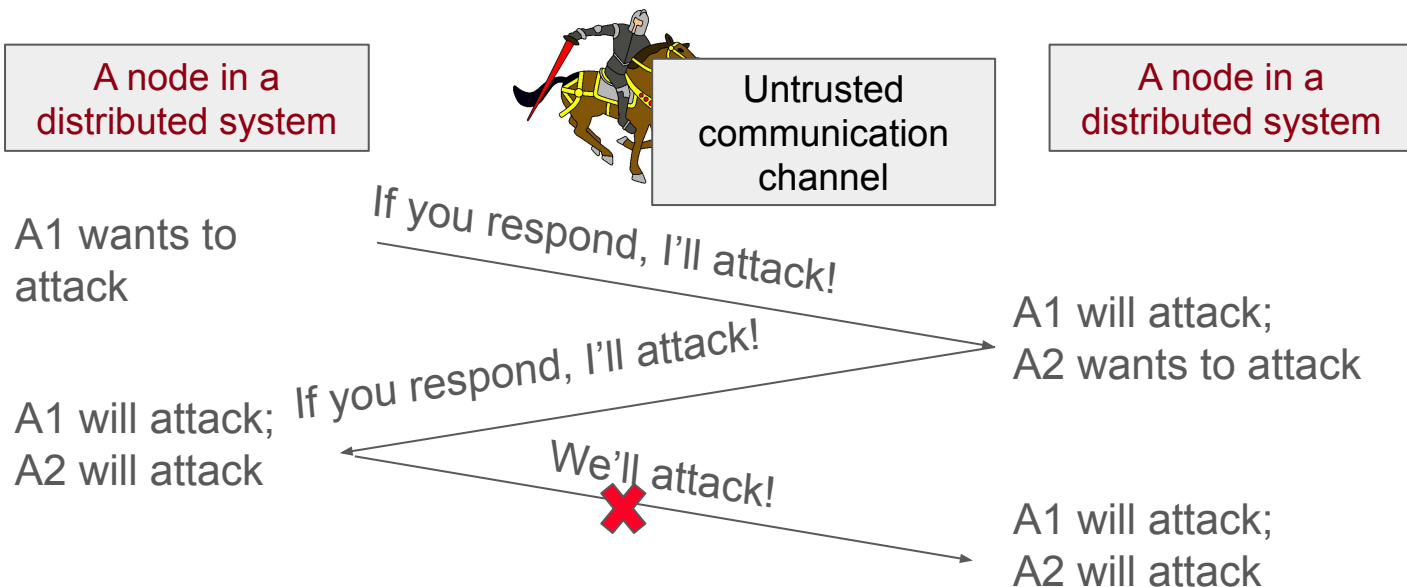
*We'll attack!*



A1 will attack;  
A2 will attack

There is no solution to the two general problem.

# Solution: TCP 3-way handshake in practice



Assumption: Failure is not 100% Byzantine.

# The Two Generals Problem

The two generals problem is a classic computer science problem that remains unsolvable.

- It remains unsolvable as there is a need to be a last acknowledgement from General A1 to A2.
- And this starts a never-ending cycle of acknowledgement as the message may get lost.
- This is why the problem is unsolvable.

# Next class: The Byzantine Generals Problem